

Variants of OCC

Dynamic Adjustment & TicToc

Dynamic Adjustment

- Based on the Forward Validation scheme i.e., validate against concurrent live transactions.
- Number of aborts is reduced by using dynamic adjustment of serialization order. Supported by the use of dynamic timestamp assignment scheme.
- Serialization order of committed transactions may be different from their commit order.

Validation scheme

- Suppose we have a validating transaction T_v and a set of active transactions T_j ($j = 1, 2, \dots, n, j \neq v$).
- There are three possible types of data conflicts which can induce serialization order between T_v and T_j :
 - $RS(T_v) \cap WS(T_j) \neq \emptyset$ (write-read conflict)
Write-read conflict can be resolved by adjusting the serialization order as $T_v \rightarrow T_j$ so that the read of T_v cannot be affected by T_j 's write. Forward adjustment.
 - $WS(T_v) \cap RS(T_j) \neq \emptyset$ (read-write conflict)
Read-write conflict can be resolved by adjusting the serialization order as $T_j \rightarrow T_v$ so that the read of T_j cannot be affected by T_v 's write. Backward adjustment.
 - $WS(T_v) \cap WS(T_j) \neq \emptyset$ (write-write conflict)
Write-write conflict can be resolved by adjusting the serialization order as $T_v \rightarrow T_j$ so that the write of T_v cannot overwrite T_j 's write. Forward adjustment.

Implementation details

- SOT (T_i)
 - To indicate the relative serialization order of the transactions.
 - Initial value = ∞ (INT_MAX)
 - If $SOT(T_i) < \infty$, it means it has been backward adjusted before a committed transaction.
- ATS (T_v)
 - When T_v comes to validation, the set of active transactions T_i such that $SOT(T_v) \geq SOT(T_i)$.
- BTS (T_v)
 - Set of active transactions T_j such that $SOT(T_j) < SOT(T_v)$
- $TR(T_v, D_p) = WTS(D_p)$ when T_v read D_p .

Validation Part 1

- First part of Validation test is used only for those validating transactions which have been backward adjusted. It is to check whether:
 - All the read operations of T_v have been read from the committed transactions T_c whose $SOT(T_c) < SOT(T_v)$
 - Whether T_v 's write is invalidated. This is done by comparing $SOT(T_v)$ with $WTS(D_p)$ and $RTS(D_p)$ of the data items D_p in T_v 's write set.

```
part-one( $T_v$ )
{
  if  $SOT(T_v) \neq \infty$  then
  {
    for  $\forall D_p \in RS(T_v)$ 
    {
      if  $TR(T_v, D_p) > SOT(T_v)$  then
        restart( $T_v$ );
    }
    for  $\forall D_p \in WS(T_v)$ 
    {
      if  $SOT(T_v) < RTS(D_p)$  or  $SOT(T_v) < WTS(D_p)$  then
        restart( $T_v$ );
    }
  }
}
```

Validation Part 2

- The purpose of part two is to detect read-write conflicts between active transactions and the validating transactions.
- Here we compare the WS (T_v) with the RS (T_i) where $T_i \in \text{ATS}(T_v)$
- The conflicting transactions T_i are added to $\text{BTlist}(T_v)$ to indicate that T_j needs to be backward adjusted before T_v .

```
part_two( $T_v$ )
{
     $\text{BTlist}(T_v) = \emptyset$  ;

    for  $\forall T_j \in \text{ATS}(T_v)$ 
    {
        for  $\forall D_p \in \text{WS}(T_v)$ 
        {
            if  $D_p \in \text{RS}(T_j)$  then
                 $\text{BTlist}(T_v) = \text{BTlist}(T_v) \cup T_j$  ;
        }
    }
}
```

Validation Part 3

- The third part of the test is to detect whether a backward-adjusted transaction T_j also needs forward adjustment w.r.t T_v
- Compares the RS (T_v) with WS (T_j) where $T_j \in \text{BTS}(T_v)$ or $\text{BTlist}(T_v)$
- Compares the WS (T_v) with WS (T_j) where $T_j \in \text{BTS}(T_v)$ or $\text{BTlist}(T_v)$
- If either one of them is non-empty, T_j has serious conflict with T_v .
- Conflict resolution can be implemented to abort transactions based on priority.
- In our code, we abort the current validating transaction T_v .

```
part_tree( $T_v$ )
{
    for  $\forall T_j \in \text{BTS}(T_v) \cup \text{BTlist}(T_v)$ 
    {
        for  $\forall D_p \in \text{RS}(T_v)$ 
        {
            if  $D_p \in \text{WS}(T_j)$  then
                conflict_resolution( $T_v, T_j$ );
        }

        for  $\forall D_p \in \text{WS}(T_v)$ 
        {
            if  $D_p \in \text{WS}(T_j)$  then
                conflict_resolution( $T_v, T_j$ );
        }
    }
}
```

Validation Part 4

- When the validating transaction reaches part 4, it is guaranteed to commit.
- The purpose is to assign a final commitment timestamp to the validating transaction and to update the necessary timestamps of the data items.

```
part_four( $T_v$ )
{
    if  $SOT(T_v) = \infty$  then
         $SOT(T_v) = validation\_time$  ;

    for  $\forall T_j \in BTlist(T_v)$ 
         $SOT(T_j) = SOT(T_v) - \epsilon$  ; //infinitesimal quantity

    for  $\forall D_p \in RS(T_v)$ 
         $RTS(D_p) = SOT(T_v)$  ;

    for  $\forall D_p \in WS(T_v)$ 
         $WTS(D_p) = SOT(T_v)$  ;

    commit  $WS(T_v)$  to database;
}
```


TicToc: Time Travelling OCC

- Traditional OCC algorithms assign static timestamps, essentially agreeing on a fixed sequential schedule -> eases conflict detection, but limits concurrency.
- TicToc does not allocate static timestamps -> does not restrict the set of potential orderings.
- Instead, a transaction's timestamp is calculated lazily at its commit time in a distributed manner based on the tuples it accesses.
 - Distributed nature -> avoids all of the bottlenecks inherent in centralized timestamp allocation schemes -> highly scalable algorithm.
 - Laziness allows the DBMS to exploit more parallelism in the workload -> reduces aborts and improves performance.

Example

- Consider the following example involving two concurrent transactions, A and B, and two tuples, x and y. The transactions invoke the following sequence of operations:
 - A read (x)
 - B write (x)
 - B commits
 - A write (y)
- If $ts(A) < ts(B)$, then A can commit since the interleaving of operations is consistent with the timestamp order.
- However if $ts(A) > ts(B)$, A must eventually abort since committing it would violate the schedule imposed by timestamp order.

Lazy Timestamp Management

- To encode the serialization information in the tuples, each data version has a valid range of timestamps bounded by the write timestamp (wts) and the read timestamp (rts) -> a particular version is created at timestamp wts and is valid until timestamp rts.
- A version read by a transaction is valid iff that transaction's commit timestamp is in between the version's wts and rts.
- A write by a transaction is valid iff the transaction's commit timestamp is greater than the rts of the previous version.

$$\begin{aligned} & \exists \text{ commit_ts}, \\ & (\forall v \in \{\text{versions read by } T\}, v.\text{wts} \leq \text{commit_ts} \leq v.\text{rts}) \\ & \wedge (\forall v \in \{\text{versions written by } T\}, v.\text{tuple.rts} < \text{commit_ts}) \end{aligned}$$

Implementation details

- A read set (RS) and write set (WS) of tuples for each transaction.
- Each entry in the RS and WS is encoded as {tuple, data, wts, rts}
 - tuple is a pointer to the tuple in the database
 - data is the data value read by the transaction
 - wts and rts are the timestamps copied from the tuple when it was accessed by the transaction.
- The value and timestamps must be loaded atomically to guarantee that the value matches timestamps.

Algorithm 1: Read Phase

Data: read set RS , tuple t

1 $r = RS.get_new_entry()$

2 $r.tuple = t$

Atomically load wts, rts, and value

3 $\langle r.value = t.value, r.wts = t.wts, r.rts = t.rts \rangle$

Validation Phase

- In the validation phase, TicTic uses the timestamps stored in the transaction's read and write sets to compute its commit timestamp.
- Then, the algorithm checks whether the tuples in the transaction's read set are valid based on this commit timestamp.
- Step 1: Lock all the tuples in the transaction's WS in their primary key order -> no concurrent updates and no deadlocks.

Algorithm 2: Validation Phase

Data: read set RS , write set WS

Step 1 – Lock Write Set

1 **for** w in $sorted(WS)$ **do**

2 $lock(w.tuple)$

3 **end**

Step 2 – Compute the Commit Timestamp

4 $commit_ts = 0$

5 **for** e in $WS \cup RS$ **do**

6 **if** e in WS **then**

7 $commit_ts = \max(commit_ts, e.tuple.rts + 1)$

8 **else**

9 $commit_ts = \max(commit_ts, e.wts)$

10 **end**

11 **end**

Step 3 – Validate the Read Set

12 **for** r in RS **do**

13 **if** $r.rts < commit_ts$ **then**

Begin atomic section

14 **if** $r.wts \neq r.tuple.wts$ **or** $(r.tuple.rts \leq commit_ts$ **and**
 $isLocked(r.tuple)$ **and** $r.tuple$ not in W) **then**

15 $abort()$

16 **else**

17 $r.tuple.rts = \max(commit_ts, r.tuple.rts)$

18 **end**

End atomic section

19 **end**

20 **end**

Validating the tuples in the Read Set

- If the transaction's `commit_ts` \leq `rts` of the read set entry, then the invariant holds and no further action.
- If the entry's `rts` $<$ `commit_ts`, it is not clear whether the local value is still valid or not at `commit_ts`.
 - If another transaction has modified the tuple at a logical time between the local `rts` and `commit_ts` \rightarrow the transaction has to abort. Otherwise, if no transaction has modified the tuple, `rts` can be extended to be greater than or equal to `commit_ts`, making the version valid at `commit_ts`. To check this, compare the local `wt`s to the latest `wt`s.
 - If `wt`s matches, but the tuple is already locked by a different transaction, it is not possible to extend the `rts` either.
 - If the `rts` is extensible or if the version is already valid at `commit_ts`, the `rts` of the tuple can be extended to at least `commit_ts`.

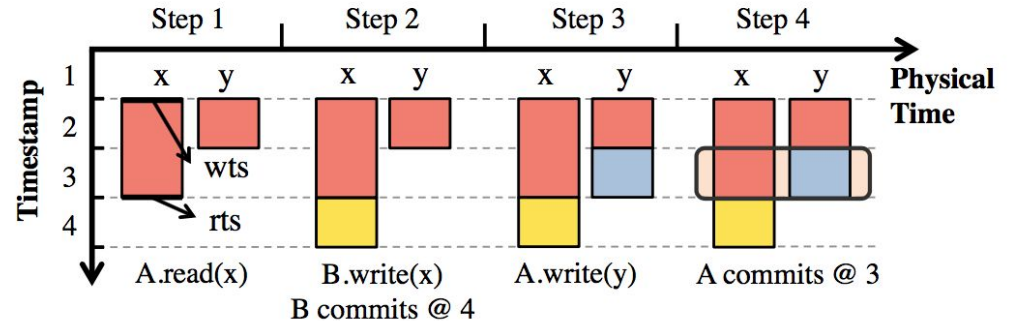
Write Phase

Example

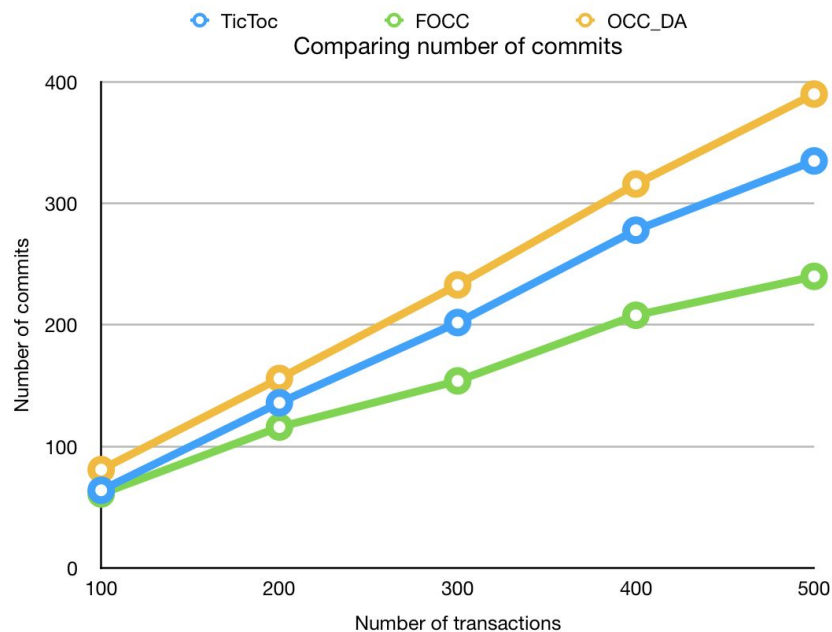
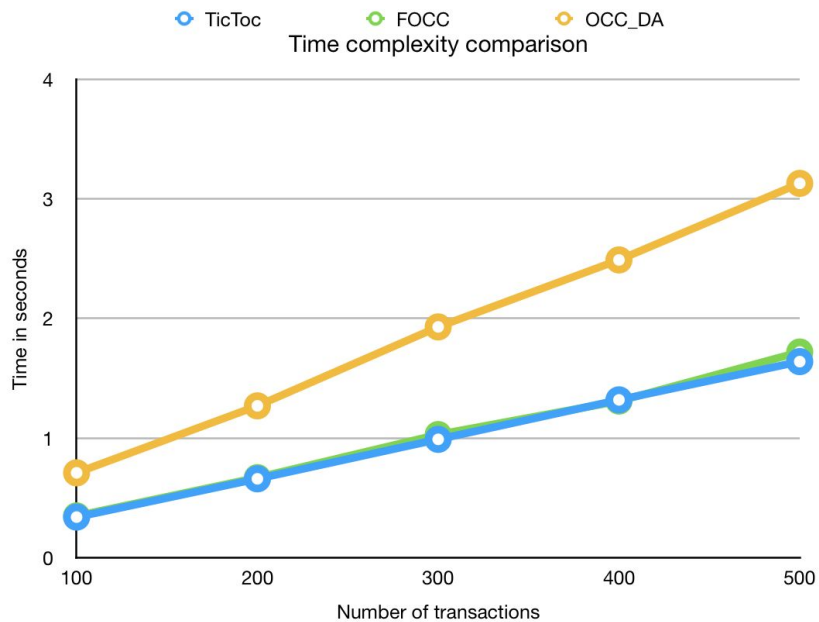
Algorithm 3: Write Phase

Data: write set WS , commit timestamp $commit_ts$

```
1 for  $w$  in  $WS$  do  
2    $write(w.tuple.value, w.value)$   
3    $w.tuple.wts = w.tuple.rts = commit\_ts$   
4    $unlock(w.tuple)$   
5 end
```



Comparison with Vanilla FOCC



Thank You!